

Lecture 4 – Key Exchange and Message Encryption

Chloe Martindale


2023

These notes are for the lecture course Cryptology at the University of Bristol. If you come across any typos, please email chloe.martindale@bristol.ac.uk.

4.1 Diffie-Hellman key exchange

In Lecture 1 we saw the *one-time pad*, which is a secret known by multiple people which then can be used for cryptography. However, having a one-time pad with which we can work requires secure offline communication, which for most real-world scenarios is not practical and definitely not cost-effective. The cryptographic solution to this is to use a *key-exchange* algorithm, which does exactly what it says on the tin: It allows two (or more but we focus on two for now) parties who communicate over an open channel to compute a shared secret value, known only to them, which they can then use to encrypt communication between them.

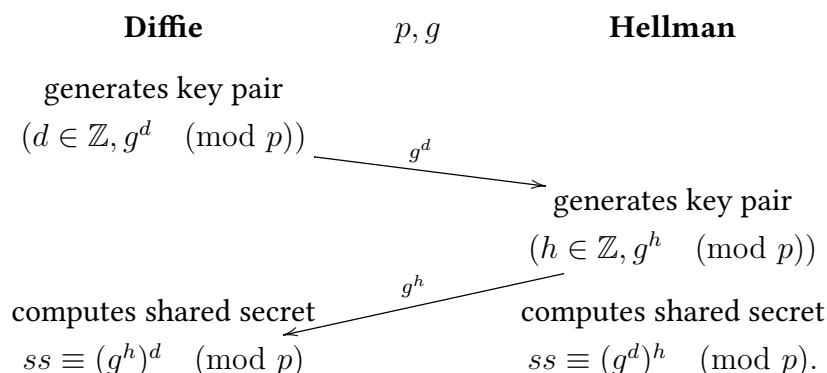
The abstract idea of a key exchange is as follows: suppose Diffie and Hellman want compute a shared secret key (read: a shared value that noone else can compute).



Here is a (overly simple) message encryption scheme built using such a key-exchange:

1. Alice and Bob compute ss via a key-exchange and encode it as a bit string.
2. Alice encodes her plaintext message m as a a bit string, computes the ciphertext $c = m \oplus s$, and sends c to Bob.
3. Bob decrypts the ciphertext via $m = c \oplus s$.

So, how do we instantiate such a key-exchange? The most basic version of the Diffie-Hellman key exchange uses exponentiation modulo a large prime p . A prime p and a nonzero element $g \pmod{p}$ is fixed in a public setup phase, and the key exchange is as follows:



In order for this to define a *cryptosystem*, we need more than just mathematical validity. We need any attack method to be much much slower than the algorithms used by the honest participants. ‘Much slower’ is something that we need to define in order to understand how to choose security parameters; for this we introduce a *security parameter* λ , which will be some smallish positive integer (often 128 in real-world scenarios), and which we use to discuss the amount of time an algorithm takes in terms of λ .

When we say that we want a computation to be ‘fast’ or *polynomial-time* we mean ‘the number of basic operations for said computation is polynomial in λ ’, i.e., you can abstractly compute an upper bound on the number of basic operations (e.g. addition) needed as a polynomial in λ . When we say that we want a computation to be ‘slow’ we mean ‘the number of basic operations is exponential or subexponential in λ ’, meaning that the number of basic operations for said computation is lower bounded by $O(2^\lambda)$ or $O(\lambda^\alpha \log_2(\lambda)^{1-\alpha})$ for some $\alpha \in (0, 1)$ respectively; these are referred to *exponential* and *subexponential* algorithms respectively. In practise, if this is true for a reasonable α (for example, for RSA which we will see below, $\alpha = 1/3$), we can increase λ to a size for which polynomial time calculations are at most milliseconds and subexponential calculations would take years.

For Diffie-Hellman, if we choose p so that $\lambda \approx \log_2(p)$, then exponentiation mod p should be easy/fast/polynomial in λ (more on this later) and the *discrete logarithm problem*, or computing d^{th} or h^{th} roots mod p , should be hard/slow/subexponential in λ (more on this later too).

4.2 ElGamal encryption

A more sophisticated method to use the Diffie-Hellman key exchange to build an encrypted messaging protocol is called *ElGamal* encryption. We’ll see next week why this a better way of doing message encryption than the “Diffie-Hellman + XOR” method. ElGamal works as follows:

- **Setup:**

1. Diffie chooses a prime p and a generator g of $\mathbb{Z}/p\mathbb{Z} - \{0\}$.
2. Diffie chooses a random secret $d \in \{1, \dots, p-1\}$ and computes $pk_D = g^d \pmod{p}$.
3. Diffie sends his public key (p, g, pk_D) to Hellman.

- **Encryption:**

1. Hellman chooses a random secret $h \in \{1, \dots, p-1\}$ and computes $pk_H = g^h \pmod{p}$.
2. Hellman computes the shared secret $ss = pk_D^h \pmod{p}$.
3. Hellman computes the encrypted message $enc_m = m \cdot ss$.
4. Hellman sends the ciphertext (pk_H, enc_m) to Diffie.

- **Decryption:**

1. Diffie computes the shared secret $ss = pk_H^d \pmod{p}$.
2. Diffie computes the ciphertext $m = enc_m \cdot ss^{-1} = enc_m \cdot pk_H^{p-1-d} \pmod{p}$.

Observations:

- Note that $ss^{-1} = pk_H^{p-1-d}$ as

$$ss \cdot pk_H^{p-1-d} = g^{dh} \cdot g^{h(p-1-d)} = (g^{p-1})^h = 1 \pmod{p}.$$

- If m is known, the shared secret ss can be recovered from the ciphertext, so use a new secret h for each message.

4.3 RSA

This section is about RSA, named after Rivest, Shamir, and Ademan. RSA was the first public key encryption (PKE) and signature system, and is still in wide use today.

The basic RSA public key encryption system consists of 3 steps: a setup phase for key generation by the user (**KeyGen**), encryption of a message m by a second party (**Encrypt**), and decryption of the message m by the user (**Decrypt**). The algorithms for these steps are below. We have coloured the users secrets in **red** and the public values in **green**.

KeyGen

1. Pick primes $p \neq q$ of bit length λ .
2. Compute $n = p \cdot q$ and $\varphi(n) = (p-1)(q-1)$.
3. Pick e coprime to $\varphi(n)$.

4. Compute $d = e^{-1} \pmod{\varphi(n)}$.

5. Generate key pair

$$\text{pk}, \text{sk} = (e, n), (d, n).$$

Encrypt

1. Pick $m \in \mathbb{Z}_{[0, n-1]}$.

2. Compute $c \equiv m^e \pmod{n}$.

3. Send c .

Decrypt

1. Compute $c^d \equiv m \pmod{n}$.

In order for this to be a valid system, there are two steps that don't obviously mathematically check out: step 4 of **KeyGen** (does the inverse exist?) and step 1 of **Decrypt**.

Recall from last week that $a \pmod{b}$ is invertible if and only if a and b are coprime, so step 4 of **KeyGen** is valid.

For the **Decrypt** step, note that if m is invertible mod n then Fermat's Little Theorem implies that $m^{\varphi(n)} \equiv 1 \pmod{n}$. Let $k \in \mathbb{Z}$ be such that $ed = 1 + k\varphi(n)$. Then

$$c^d \equiv (m^e)^d \equiv m^{1+k\varphi(n)} \equiv m \cdot (m^{\varphi(n)})^k \equiv m \cdot 1^k \equiv m \pmod{n},$$

so the decrypt step is valid (if m is invertible mod n , actually also if it's not but that requires some more steps).

For RSA to work as a cryptosystem:

- Step 2 of **KeyGen** needs to be fast, i.e., we need to be able to multiply fast.
- Step 4 of **KeyGen** needs to be fast, i.e., we need to be able to compute inverses mod $\varphi(n)$ fast.
- In Step 5 of **KeyGen**, an attacker shouldn't be able to compute d from (e, n) . Because Step 4 is fast, if the attacker knows $\varphi(n)$ then they can compute d fast. So computing $\varphi(n)$ from n should be slow. As $\varphi(n)$ is easy to compute from p and q , factoring n should be slow.
- Step 2 of **Encrypt** needs to be fast, i.e., we need to be able to exponentiate mod n fast.
- An attacker should also not be able to recover m from c , so computing e^{th} roots mod n should be slow.

4.4 Fast multiplication, exponentiation, and inversion

Let us first focus on the computations we want to be fast in RSA. To multiply fast, we use a method called ‘double-and-add’. To see how this works let’s consider how we would multiply p by q in Step 2 of **KeyGen**. We first write q in binary as (q_λ, \dots, q_0) , or in other words

$$q = \sum_{i=0}^{\lambda} q_i 2^i,$$

where $q_i \in \{0, 1\}$. In particular

$$pq = \sum_{i=0}^{\lambda} q_i \cdot (2^i p).$$

We can compute the $2^i p$ terms just by repeated doubling – each doubling is just one addition (which is a basic operation) so this is very efficient.

- **Double:** Compute

$$\begin{aligned} 2^0 \cdot p &= p \rightarrow 0 \text{ additions} \\ 2^1 \cdot p &= p + p \rightarrow 1 \text{ addition} \\ 2^2 \cdot p &= 2p + 2p \rightarrow 1 \text{ addition} \\ &\dots \\ 2^\lambda \cdot p &= 2^{\lambda-1}p + 2^{\lambda-1}p \rightarrow 1 \text{ addition.} \end{aligned}$$

The doubling step costs λ additions, so is polynomial in λ i.e. fast.

Now to get $p \cdot q$ we just have to add together the terms $2^i \cdot p$ together for which $q_i = 1$. So, let q_{i_0}, \dots, q_{i_k} be the non zero coefficients of the binary expansion of q .

- **Add:** Compute

$$p \cdot q = (2^{i_0} \cdot p) + \dots + (2^{i_k} \cdot p).$$

The adding step costs $k \leq \lambda$ additions.

In total, the double-and-add method costs at most 2λ basic operations, so is ‘fast’.

To exponentiate mod n fast, we play the same game. To see how this works let’s consider computing $m^e \pmod n$ as we do in **Encrypt**. This time we use the binary expansions of $e = (e_\lambda, \dots, e_0)$, so in other words

$$e = \sum_{i=0}^{\lambda} e_i 2^i,$$

where $e_i \in \{0, 1\}$. In particular

$$m^e = m^{\sum_{i=0}^{\lambda} e_i 2^i} = \prod_{i=0}^{\lambda} (m^{2^i})^{e_i}.$$

We can compute the $m^{2^i} \pmod n$ terms just by repeated squaring – each squaring is at most one multiplication (maybe you get some cancellation so it could be less), each of which we just saw is at most 2λ basic operations.

- **Square:** Compute

$$\begin{aligned} m^{2^0} &\equiv m \pmod n \rightarrow 0 \text{ squarings} \\ m^{2^1} &\equiv m^2 \pmod n \rightarrow 1 \text{ squaring} \\ m^{2^2} &\equiv (m^2)^2 \pmod n \rightarrow 1 \text{ squaring} \\ &\dots \\ m^{2^\lambda} &\equiv (m^{2^{\lambda-1}})^2 \pmod n \rightarrow 1 \text{ squaring.} \end{aligned}$$

The squaring step costs λ squarings, so is at most $2\lambda^2$ basic operations.

Note that the fact that we are doing the computations $\pmod n$ here is very important – for large λ you would quickly run into memory problems otherwise. All that remains now to get $m^e \pmod n$ is to multiply together the terms $m^{2^i} \pmod n$ together for which $e_i = 1$. So, let e_{i_0}, \dots, e_{i_k} be the non zero coefficients of the binary expansion of e .

- **Multiply:** Compute

$$m^e \pmod n \equiv m^{2^{i_0}} \times \dots \times m^{2^{i_k}} \pmod n.$$

The multiplication step then costs $k \leq \lambda$ multiplications, so $\leq 2\lambda^2$ basic operations.

From these calculations, we see that square-and-multiply can always be performed in $\leq 4\lambda^2$ basic operations, so is polynomial time, i.e. ‘fast’. In practise you can do quite a bit better! But we won’t go into that now.

If we look now at our list of computations that we want to be fast for RSA to work as a cryptosystem, we’ve tackled multiplication and exponentiation, and only inversion $\pmod{\varphi(n)}$ remains.

We saw two algorithms for computing inverses last week: Euclid’s corollary and Fermat’s Little Theorem. However, as $\varphi(n)$ may have many small factors the most efficient algorithm here would be to combine Euclid’s corollary with Sun-Tzu’s Remainder Theorem—see the exercise sheet for this week.

4.5 The Discrete Logarithm Problem

Just as with RSA, the validity of ElGamal encryption and more generally the Diffie-Hellman key exchange relies on certain computations being easy (fast) or hard (slow). For Diffie-Hellman, exponentiation mod p should be easy/fast/polynomial-time, which we have already seen it can be with square-and-multiply.

The fundamental problem that should be hard/slow/(sub)exponential-time for Diffie-Hellman is the *discrete logarithm problem*, that is, computing $d \in [0, p-1]$ given $g \pmod{p}$ and $g^d \pmod{p}$.¹

There are instances where this might be very easy, for example if $g = p-1 \equiv -1 \pmod{p}$ then the only values of g^d or g^h that can occur are -1 and 1 so finding a root is very easy. To avoid this, we want g^d to be able to take as many values as possible.

For the Diffie-Hellman key exchange, our key space is $(\mathbb{Z}/p\mathbb{Z})^*$, where p is a prime, which we saw last week is defined by

$$(\mathbb{Z}/p\mathbb{Z})^* = \{1 \pmod{p}, \dots, p-1 \pmod{p}\};$$

we saw also that $(\mathbb{Z}/p\mathbb{Z})^*$ is a cyclic group under multiplication \pmod{p} . That is, there exists a $g \pmod{p}$ such that

$$(\mathbb{Z}/p\mathbb{Z})^* = \{g \pmod{p}, g^2 \pmod{p}, \dots, g^{p-1} \pmod{p}\}.$$

The reason that it is important for Diffie-Hellman that $(\mathbb{Z}/p\mathbb{Z})^*$ is cyclic is because it is possible to choose $g \pmod{p} \in (\mathbb{Z}/p\mathbb{Z})^*$ for which $g^d \pmod{p}$ takes $p-1$ different values: so that there is exactly one valid private key (d) for any given public key (g^d). In fact, the reason that we chose the letter ‘ g ’ when setting up the key exchange is because we typically choose a *generator* for the cyclic group $(\mathbb{Z}/p\mathbb{Z})^*$ (or a subgroup of that, but for simplicity we ignore that for now).

In conclusion, for our Diffie-Hellman setup, we choose p prime and g a generator of the group $(\mathbb{Z}/p\mathbb{Z})^*$. Observe that this choice only avoids the most obvious reason for the discrete logarithm problem (computing d from $g^d \pmod{p}$) being easy; we’ll get to other algorithms to compute discrete logarithms in due course.

4.6 SRT vs. the Discrete Logarithm Problem

Later in the course we will look at some of the best known classical (that is ‘not quantum’) algorithms to attack this problem, so computing $d \in [0, p-1]$ given $g^d \pmod{p}$. In some contexts though we already have the tools we need: Sun-Tzu’s Remainder Theorem!

We define the *order* d of an element g of a group G with group operation $*$ and identity id as the minimum positive integer d such that $\underbrace{g * \dots * g}_{d \text{ times}} = id$. Of course in our context

this looks like the minimum positive integer d such that $g^d \equiv 1 \pmod{p}$. In particular, the generator g that we’ve been using in Diffie-Hellman and ElGamal has order $p-1$ (if you don’t immediately see why, look back at Fermat’s Little Theorem and take some time to think about this).

Going back to the example with $p = 7$, you can hopefully now spot that $3^2 \equiv 2 \pmod{7}$ is an element of order 3, and $3^3 \equiv 6 \pmod{7}$ is an element of order 2. These are examples

¹Computing d given g^d if $g \in \mathbb{R}$ is something you’ve seen before: namely computing logarithms base g . However, the problem turns out to be fundamentally different when instead of working in a continuous solution set like \mathbb{R} we are working in a discrete solution set like $\mathbb{Z}/p\mathbb{Z}$ —hence the name the *Discrete Logarithm Problem*.

of a more general concept: if g generates $\mathbb{Z}/p\mathbb{Z} - \{0\}$ (so has order $p - 1$) and ℓ divides $p - 1$, then $g^{\frac{p-1}{\ell}} \pmod{p}$ has order ℓ (exercise: prove this); this gives an easy way of finding elements of a given order—this is going to be very useful in the following example.

Example Suppose that we want to solve the following Discrete Logarithm Problem: find $a \in \mathbb{Z}$ such that $2^a \equiv 17 \pmod{37}$, and you are given that 2 is a generator of the multiplicative group $(\mathbb{Z}/37\mathbb{Z})^*$. Then as a is in the exponent, it suffices to compute $a \pmod{36}$ (because of Fermat's Little Theorem). If we want to compute $a \pmod{36}$, by Sun-Tzu's Remainder Theorem it suffices to compute $a \pmod{4}$ and $a \pmod{9}$. This is something we can just do by brute force and observation, but there is a more efficient way using the group theory above:

- To compute $a \pmod{4}$, we first compute $a \pmod{2}$. Write $a = a_0 + 2a_1$ where $a_0 \in \{0, 1\}$. By the above, we know that $2^{(p-1)/2} = 2^{18}$ is an element of order 2. Substituting for a , a_0 , and a_1 , we get the following equalities mod 37

$$-1 \equiv 17^{18} \equiv (2^a)^{18} \equiv 2^{18a_0 + 36a_1} \equiv (2^{18})^{a_0} \cdot (2^{36})^{a_1} \equiv (-1)^{a_0},$$

from which we can read off that $a_0 = 1$, so $a \equiv 1 \pmod{2}$.

- Now we compute $a \pmod{4}$. We know that $a \equiv 1 \pmod{2}$, so there exist a_1, a_2 with $a_1 \in \{0, 1\}$ such that $a = 1 + 2a_1 + 4a_2$. By the above, we know that $2^{(p-1)/4} = 2^9$ is an element of order 4. Substituting for a , a_1 , a_2 , we get the following equalities mod 37

$$31 \equiv 17^9 \equiv (2^a)^9 \equiv 2^{9 + 18a_1 + 36a_2} \equiv 2^9 \cdot (2^{18})^{a_1} \cdot (2^{36})^{a_2} \equiv 6 \cdot (-1)^{a_1},$$

from which we can read off that $a_1 = 1$, so $a \equiv 3 \pmod{4}$.

- To compute $a \pmod{9}$, we first compute $a \pmod{3}$. Write $a = a_0 + 3a_1$ where $a_0 \in \{0, 1, 2\}$. By the above, we know that $2^{(p-1)/3} = 2^{12}$ is an element of order 3. Substituting for a , a_0 , a_1 we get the following equations mod 37

$$26 \equiv 17^{12} \equiv (2^a)^{12} \equiv 2^{12a_0 + 36a_1} \equiv (2^{12})^{a_0} \cdot (2^{36})^{a_1} \equiv 26^{a_0},$$

from which we can read off that $a_0 = 1$.

- Now we compute $a \pmod{9}$. We know that $a \equiv 1 \pmod{3}$, so there exist a_1, a_2 with $a_1 \in \{0, 1, 2\}$ such that $a = 1 + 3a_1 + 9a_2$. By the above, we have that $2^{(p-1)/9} = 2^4$ is an element of order 9. Substituting for a , a_1 , a_2 , we get the following equations mod 37

$$12 \equiv 17^4 \equiv (2^a)^4 \equiv 2^{4 + 12a_1 + 36a_2} \equiv 2^4 \cdot (2^{12})^{a_1} \cdot (2^{36})^{a_2} \equiv 16 \cdot 26^{a_1},$$

from which we can read off that $a_1 = 2$, so $a \equiv 7 \pmod{9}$.

- Now we know that $a \equiv 3 \pmod{4}$ and $a \equiv 7 \pmod{9}$, which by CRT we know uniquely defines $a \pmod{36}$. We can compute this via Euclid's algorithm as we've done before, giving $a \equiv 7 \pmod{36}$.

The above method is in no way specific to the numbers we have chosen (37, 4, 9, etc): it is in fact an example of an algorithm that works in general. This trick of using Sun-Tzu's Remainder Theorem to attack the Discrete Logarithm Problem is due to Pohlig and Hellman and is therefore referred to as the Pohlig-Hellman algorithm.